2024 was also an outstanding year for our data center AI business.

We launched our AMD Instinct MI325X GPUs with industry-leading inferencing performance, accelerated our AI hardware roadmap to deliver an annual cadence of new accelerators, significantly enhanced our ROCm software suite and strengthened our partnerships with many of the world's largest AI providers.

Leading cloud providers and AI innovators adopted AMD Instinct MI300 accelerators at scale to power their AI workloads. Microsoft highlighted the TCO advantages of AMD Instinct MI300X accelerators when running multiple Microsoft Copilot services that use the latest GPT-4 models. Meta announced broad deployment of MI300X for inference, including serving all live traffic for its most demanding Llama 405B frontier model exclusively on AMD GPUs.

Microsoft, Oracle Cloud, IBM Cloud and more than a dozen other AI-specialized cloud providers adopted AMD Instinct accelerators to power their public cloud instances, including flagship instances available on Microsoft Azure that scale up to thousands of GPUs for AI inference and training and high-performance computing workloads.

Annual Cadence of New AI Accelerators

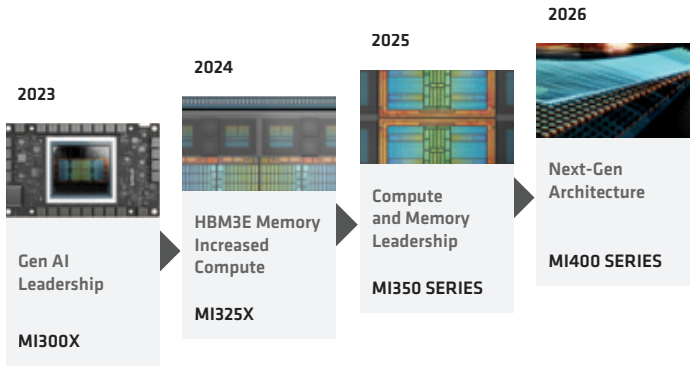

The El Capitan supercomputer at Lawrence Livermore National Laboratory, powered by AMD Instinct MI300A APUs, debuted as the fastest computer in the world on the latest TOP500 list. AMD proudly powers two of the top three exascale computers in the United States, and El Capitan is one of them. Adoption of our AMD EPYC and AMD Instinct processors with the broader supercomputer customer base also grew in the year, with AMD-based systems comprising 50% of the top ten fastest and 40% of the ten most energy-efficient supercomputers in the world.